

```

# Import necessary libraries
from pyspark.sql import SparkSession
from pyspark.sql.functions import *
import pyspark.sql.functions as F
from pyspark.ml.stat import Correlation

from tqdm import tqdm # For progress bars
from google.colab import drive
drive.mount('/content/drive')
# Data manipulation and visualization
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Initialize Spark session
spark = SparkSession.builder \
    .appName("HomeCredit_EDA") \
    .config("spark.driver.memory", "12g") \
    .config("spark.executor.memory", "10g") \
    .config("spark.memory.fraction", "0.8") \
    .config("spark.memory.storageFraction", "0.2") \
    .config("spark.driver.maxResultSize", "2g") \
    .config("spark.sql.shuffle.partitions", "12") \
    .config("spark.default.parallelism", "12") \
    .getOrCreate()

Mounted at /content/drive

```

1. Importing database

```

# List of all data files to import
filename = [
    'application_train',    # Main training data with target variable
    'bureau',              # Credit bureau data
    'bureau_balance',      # Monthly bureau balance data
    'credit_card_balance', # Credit card monthly balance
    'installments_payments', # Payment history data
    'POS_CASH_balance',    # Point of sale and cash loan data
    'previous_application' # Previous application data
]

# Dictionary to store all dataframes
data = {}

# Import all files into dictionary
for file in tqdm(filename):
    if file in data.keys():

```

```

        print(f'{file} already exists in dictionary')
    else:
        filepath = '/content/drive/MyDrive/Home_Credit/'+f'{file}.csv'
        df = spark.read.csv(filepath, header=True, sep=',',
inferSchema=True)
        data[file] = df
100%|██████████| 7/7 [01:41<00:00, 14.55s/it]

```

1.1. Summarizing dataframes

```

"""
Create a summary table showing number of columns and observations in
each dataframe
"""

summary_results = []
for key in data.keys():
    summary_results.append([key, len(data[key].columns),
data[key].count()])
dataframe_summary = pd.DataFrame(summary_results,
columns=['Dataframe', 'Number of columns', 'Number of observations'])
dataframe_summary

{"summary":{"\n  \"name\": \"dataframe_summary\",\n  \"rows\": 7,\n  \"fields\": [\n    {\n      \"column\": \"Dataframe\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 7,\n        \"samples\": [\n          \"application_train\",\n          \"bureau\",\n          \"POS_CASH_balance\",\n          \"semantic_type\": \"\",\n          \"description\": \"\",\n          \"column\": \"Number of columns\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 41,\n            \"min\": 3,\n            \"max\": 122,\n            \"num_unique_values\": 6,\n            \"samples\": [\n              122,\n              17,\n              37\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\",\n            \"column\": \"Number of observations\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 9687522,\n              \"min\": 307511,\n              \"max\": 27299925,\n              \"num_unique_values\": 7,\n              \"samples\": [\n                307511,\n                1716428,\n                10001358\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\",\n              \"column\": \"\n            ]\n          },\n          \"type\": \"dataframe\", \"variable_name\": \"dataframe_summary\"}

```

1.2. Missing values counting

```

def analyze_missing_data(dataframe_name):
    """
    Analyze missing values in a dataframe

```

Args:

dataframe_name: Name of the dataframe in the data dictionary

Returns:

DataFrame with missing value counts and percentages

```
"""
df = data[dataframe_name]
# Count nulls in each column
missing_counts = df.select([count(when(col(c).isNull(),
c)).alias(c) for c in df.columns])
# Add a row identifier and transpose
missing_df = missing_counts.withColumn('Column',
lit('Total')).toPandas().set_index('Column').transpose()
# Calculate percentage
missing_df['Percent'] =
(missing_df['Total']*100/df.count()).round(2)
return missing_df

# Analyze missing values in each dataframe
print('Top 20 variables with missing values in each dataframe:')
for df_name in filename:
    missing_df = analyze_missing_data(df_name)
    print(f'\n{df_name.upper()} with total {data[df_name].count()}
observations:')
    print(missing_df.sort_values('Total', ascending=False).head(20))
```

Top 20 variables with missing values in each dataframe:

APPLICATION_TRAIN with total 307511 observations:

Column	Total	Percent
COMMONAREA_AVG	214865	69.87
COMMONAREA_MODE	214865	69.87
COMMONAREA_MEDI	214865	69.87
NONLIVINGAPARTMENTS_MEDI	213514	69.43
NONLIVINGAPARTMENTS_MODE	213514	69.43
NONLIVINGAPARTMENTS_AVG	213514	69.43
FONDKAPREMONT_MODE	210295	68.39
LIVINGAPARTMENTS_AVG	210199	68.35
LIVINGAPARTMENTS_MEDI	210199	68.35
LIVINGAPARTMENTS_MODE	210199	68.35
FLOORSMIN_MODE	208642	67.85
FLOORSMIN_AVG	208642	67.85
FLOORSMIN_MEDI	208642	67.85
YEARS_BUILD_AVG	204488	66.50
YEARS_BUILD_MODE	204488	66.50
YEARS_BUILD_MEDI	204488	66.50
OWN_CAR_AGE	202929	65.99
LANDAREA_MEDI	182590	59.38
LANDAREA_AVG	182590	59.38
LANDAREA_MODE	182590	59.38

BUREAU with total 1716428 observations:

Column	Total	Percent
AMT_ANNUITY	1226791	71.47
AMT_CREDIT_MAX_OVERDUE	1124488	65.51
DAYS_ENDDATE_FACT	633653	36.92
AMT_CREDIT_SUM_LIMIT	591780	34.48
AMT_CREDIT_SUM_DEBT	257669	15.01
DAYS_CREDIT_ENDDATE	105553	6.15
AMT_CREDIT_SUM	13	0.00
SK_ID_CURR	0	0.00
SK_ID_BUREAU	0	0.00
CREDIT_DAY_OVERDUE	0	0.00
CREDIT_ACTIVE	0	0.00
CREDIT_CURRENCY	0	0.00
DAYS_CREDIT	0	0.00
CNT_CREDIT_PROLONG	0	0.00
AMT_CREDIT_SUM_OVERDUE	0	0.00
CREDIT_TYPE	0	0.00
DAYS_CREDIT_UPDATE	0	0.00

BUREAU_BALANCE with total 27299925 observations:

Column	Total	Percent
SK_ID_BUREAU	0	0.0
MONTHS_BALANCE	0	0.0
STATUS	0	0.0

CREDIT_CARD_BALANCE with total 3840312 observations:

Column	Total	Percent
AMT_PAYMENT_CURRENT	767988	20.00
CNT_DRAWINGS_POS_CURRENT	749816	19.52
AMT_DRAWINGS_ATM_CURRENT	749816	19.52
CNT_DRAWINGS_ATM_CURRENT	749816	19.52
AMT_DRAWINGS_POS_CURRENT	749816	19.52
AMT_DRAWINGS_OTHER_CURRENT	749816	19.52
CNT_DRAWINGS_OTHER_CURRENT	749816	19.52
CNT_INSTALMENT_MATURE_CUM	305236	7.95
AMT_INST_MIN_REGULARITY	305236	7.95
AMT_DRAWINGS_CURRENT	0	0.00
AMT_CREDIT_LIMIT_ACTUAL	0	0.00
SK_ID_PREV	0	0.00
SK_ID_CURR	0	0.00
AMT_BALANCE	0	0.00
MONTHS_BALANCE	0	0.00
AMT_TOTAL_RECEIVABLE	0	0.00
AMT_RECIVABLE	0	0.00
AMT_RECEIVABLE_PRINCIPAL	0	0.00
AMT_PAYMENT_TOTAL_CURRENT	0	0.00
CNT_DRAWINGS_CURRENT	0	0.00

INSTALLMENTS_PAYMENTS with total 13605401 observations:

Column	Total	Percent
AMT_PAYMENT	2905	0.02
DAYS_ENTRY_PAYMENT	2905	0.02
SK_ID_PREV	0	0.00
SK_ID_CURR	0	0.00
NUM_INSTALLMENT_NUMBER	0	0.00
NUM_INSTALLMENT_VERSION	0	0.00
DAYS_INSTALLMENT	0	0.00
AMT_INSTALLMENT	0	0.00

POS_CASH_BALANCE with total 10001358 observations:

Column	Total	Percent
CNT_INSTALLMENT_FUTURE	26087	0.26
CNT_INSTALLMENT	26071	0.26
SK_ID_CURR	0	0.00
SK_ID_PREV	0	0.00
MONTHS_BALANCE	0	0.00
NAME_CONTRACT_STATUS	0	0.00
SK_DPD	0	0.00
SK_DPD_DEF	0	0.00

PREVIOUS_APPLICATION with total 1670214 observations:

Column	Total	Percent
RATE_INTEREST_PRIVILEGED	1664263	99.64
RATE_INTEREST_PRIMARY	1664263	99.64
AMT_DOWN_PAYMENT	895844	53.64
RATE_DOWN_PAYMENT	895844	53.64
NAME_TYPE_SUITE	820405	49.12
DAYS_TERMINATION	673065	40.30
DAYS_FIRST_DRAWING	673065	40.30
DAYS_FIRST_DUE	673065	40.30
DAYS_LAST_DUE_1ST_VERSION	673065	40.30
DAYS_LAST_DUE	673065	40.30
NFLAG_INSURED_ON_APPROVAL	673065	40.30
AMT_GOODS_PRICE	385515	23.08
AMT_ANNUITY	372235	22.29
CNT_PAYMENT	372230	22.29
PRODUCT_COMBINATION	346	0.02
AMT_CREDIT	1	0.00
SK_ID_PREV	0	0.00
AMT_APPLICATION	0	0.00
NAME_CONTRACT_TYPE	0	0.00
SK_ID_CURR	0	0.00

2. RELATIONSHIP BETWEEN DATAFRAMES

- **SK_ID_CURR** connects the dataframes *application_train/test*, *bureau*, *previous_application*, *POS_CASH_balance*, *installments_payment* and *credit_card_balance*
- **SK_ID_PREV** connects the dataframes *previous_application*, *POS_CASH_balance*, *installments_payment* and *credit_card_balance*
- **SK_ID_PREV** connects the dataframes *bureau* and *bureau_balance*

```
# Display relationship information
relationship_table = []
for df_name in filename:
    # Count distinct SK_ID_CURR
    if 'SK_ID_CURR' in data[df_name].columns:
        distinct_curr =
data[df_name].select(F.countDistinct("SK_ID_CURR")).collect()[0][0]
    else:
        distinct_curr = 0

    # Count distinct SK_ID_PREV
    if 'SK_ID_PREV' in data[df_name].columns:
        distinct_prev =
data[df_name].select(F.countDistinct("SK_ID_PREV")).collect()[0][0]
    else:
        distinct_prev = 0

    relationship_table.append([df_name, data[df_name].count(),
distinct_curr, distinct_prev])

# Create relationship summary table
relationship_summary = pd.DataFrame(relationship_table,
                                   columns=['Dataframe', 'Total_Rows',
'Distinct_SK_ID_CURR', 'Distinct_SK_ID_PREV'])
relationship_summary

{"summary":{"\n  \"name\": \"relationship_summary\",\n  \"rows\": 7,\n  \"fields\": [\n    {\n      \"column\": \"Dataframe\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 7,\n        \"samples\": [\n          \"application_train\",\n          \"bureau\",\n          \"POS_CASH_balance\",\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"Total_Rows\",\n        \"properties\": {\n          \"dtype\": \"number\",\n          \"std\": 9687522,\n          \"min\": 307511,\n          \"max\": 27299925,\n          \"num_unique_values\": 7,\n          \"samples\": [\n            307511,\n            1716428,\n            10001358\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      }\n    ],\n    {\n      \"column\": \"Distinct_SK_ID_CURR\",
```

```

\"Distinct_SK_ID_CURR\",\\n        \\\"properties\\\": {\\n            \\\"dtype\\\":
\\\"number\\\",\\n            \\\"std\\\": 137754,\\n            \\\"min\\\": 0,\\n
\\\"max\\\": 339587,\\n            \\\"num_unique_values\\\": 7,\\n
\\\"samples\\\": [\\n                307511,\\n                305811,\\n
337252\\n            ],\\n            \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n        }\\n    },\\n    {\\n        \\\"column\\\":
\\\"Distinct_SK_ID_PREV\\\",\\n        \\\"properties\\\": {\\n            \\\"dtype\\\":
\\\"number\\\",\\n            \\\"std\\\": 671800,\\n            \\\"min\\\": 0,\\n
\\\"max\\\": 1670214,\\n            \\\"num_unique_values\\\": 5,\\n
\\\"samples\\\": [\\n                104307,\\n                1670214,\\n
997752\\n            ],\\n            \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n        }\\n    }\\n ]\\n
n}\\\", \"type\": \"dataframe\", \"variable_name\": \"relationship_summary\"}

```

VISUALIZATION FUNCTIONS

```

def plot_dist_TARGET(data_frame, variables, num_rows):
    """
    Plot distribution of variables by TARGET variable (default vs non-
    default)

    Args:
        data_frame: The dataframe containing the data
        variables: List of variables to plot
        num_rows: Number of rows for the visualization grid
    """
    i = 0
    # Fill NA values with 0 for the selected variables
    temp_df = data_frame.select(['TARGET'] + variables).na.fill(0)
    # Split data by TARGET
    default_loans = temp_df.filter(col('TARGET') == 1).toPandas()
    non_default_loans = temp_df.filter(col('TARGET') == 0).toPandas()

    # Set up plot style
    sns.set_style('whitegrid')
    fig, axs = plt.subplots(num_rows, 2, figsize=(15, 4*num_rows))
    axs = axs.flatten() if num_rows > 1 else [axs[0], axs[1]]

    # Plot each variable
    for feature in variables:
        # Calculate correlation with TARGET
        correlation = temp_df.stat.corr('TARGET', feature)
        # Calculate median values for each group
        default_median = temp_df.filter(col('TARGET') ==
1).approxQuantile(feature, [0.5], 0)
        non_default_median = temp_df.filter(col('TARGET') ==
0).approxQuantile(feature, [0.5], 0)

        # Print statistics

```

```

        print(f'Correlation between {feature} and TARGET:
{correlation:.4f}')
        print(f'Median value for default loans:
{default_median[0]:.4f}')
        print(f'Median value for non-default loans:
{non_default_median[0]:.4f}')
        print('\n')

        # Create subplot
        ax = axs[i]
        # Modern KDE plot syntax
        sns.kdeplot(data=default_loans, x=feature, ax=ax,
label="TARGET = 1", fill=True, alpha=0.3)
        sns.kdeplot(data=non_default_loans, x=feature, ax=ax,
label="TARGET = 0", fill=True, alpha=0.3)
        ax.set_ylabel('Density', fontsize=12)
        ax.set_xlabel(feature, fontsize=12)
        ax.tick_params(axis='both', which='major', labelsize=12)
        ax.legend()
        i += 1

plt.tight_layout()
plt.show()

def plot_freq(data_frame, feature, horizontal_layout=True,
label_rotation=False):
    """
    Plot frequency count and TARGET rate for a categorical variable

    Args:
        data_frame: The dataframe containing the data
        feature: The categorical feature to plot
        horizontal_layout: Whether to arrange plots horizontally or
vertically
        label_rotation: Whether to rotate x-axis labels
    """
    # Count frequency of each category
    freq_table =
data_frame.where(col(feature).isNotNull()).groupby(feature).count().so
rt(feature, ascending=True)
    freq_table = freq_table.toPandas().rename(columns={'count':
'Number_of_contracts'})

    # Calculate default rate for each category
    default_rate =
data_frame.where(col(feature).isNotNull()).select([feature,
'TARGET']).groupBy(feature).mean()
    default_rate = default_rate.select([feature,
'avg(TARGET)']).sort(feature, ascending=True).toPandas()

```



```

# Create plots
if horizontal_layout:
    fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(15, 6))
else:
    fig, (ax1, ax2) = plt.subplots(nrows=2, figsize=(12, 10))
    plt.subplots_adjust(hspace=0.4)

# Plot frequency - modern Seaborn syntax
sns.barplot(x=feature, y="Number_of_contracts", data=freq_table,
ax=ax1)
ax1.set_title(f'Frequency of {feature}', fontsize=14)
if label_rotation:
    plt.setp(ax1.get_xticklabels(), rotation=90, ha='right')

# Plot default rate - modern Seaborn syntax
sns.barplot(x=feature, y='avg(TARGET)', data=default_rate, ax=ax2)
ax2.set_title(f'Default Rate by {feature}', fontsize=14)
ax2.set_ylabel('Default Rate', fontsize=12)
ax2.set_ylim(0, default_rate['avg(TARGET)'].max() * 1.1) # Add
some headroom
if label_rotation:
    plt.setp(ax2.get_xticklabels(), rotation=90, ha='right')

plt.tight_layout()
plt.show()

def plot_dist(data_frame, feature, color):
    """
    Plot distribution of a numeric variable

    Args:
        data_frame: The dataframe containing the data
        feature: The numeric feature to plot
        color: Color for the plot
    """
    # Extract and clean data
    feature_data = data_frame.select(feature).toPandas().dropna()

    # Create figure
    plt.figure(figsize=(12, 6))

    # Main histogram with KDE
    sns.histplot(
        data=feature_data,
        x=feature,
        kde=True,
        color=color,
        alpha=0.7,
        stat="density",
        bins=30
    )

```

```

)

# Add descriptive statistics as text
stats = data_frame.select(
    F.mean(feature).alias('mean'),
    F.min(feature).alias('min'),
    F.max(feature).alias('max'),
    F.stddev(feature).alias('std')
).collect()[0]

stats_text = (
    f"Mean: {stats['mean']:.2f}\n"
    f"Min: {stats['min']:.2f}\n"
    f"Max: {stats['max']:.2f}\n"
    f"Std: {stats['std']:.2f}"
)

# Position stats text in the upper right of the plot
plt.annotate(
    stats_text,
    xy=(0.95, 0.95),
    xycoords='axes fraction',
    ha='right',
    va='top',
    bbox=dict(boxstyle="round,pad=0.5", fc="white", alpha=0.8)
)

plt.title(f"Distribution of {feature}", fontsize=15)
plt.xlabel(feature, fontsize=12)
plt.ylabel("Density", fontsize=12)
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

```

2.1. Dataframe `installments_payments`

```

key = 'installments_payments'
print(f'Dataframe {key} includes {len(data[key].columns)} features and {data[key].count()} observations.')
data[key].printSchema()

```

Dataframe `installments_payments` includes 8 features and 13605401 observations.

```

root
|-- SK_ID_PREV: integer (nullable = true)
|-- SK_ID_CURR: integer (nullable = true)
|-- NUM_INSTALLMENT_VERSION: double (nullable = true)
|-- NUM_INSTALLMENT_NUMBER: integer (nullable = true)
|-- DAYS_INSTALLMENT: double (nullable = true)
|-- DAYS_ENTRY_PAYMENT: double (nullable = true)

```

```

|-- AMT_INSTALLMENT: double (nullable = true)
|-- AMT_PAYMENT: double (nullable = true)

# Create derived features
installments_df = data[key]

# Calculate amount left (difference between prescribed and actual
payment)
installments_df = installments_df.withColumn('AMT_LEFT',
col('AMT_INSTALLMENT') - col('AMT_PAYMENT'))

# Calculate days late
installments_df = installments_df.withColumn('DAYS_LATE',
col('DAYS_INSTALLMENT') - col('DAYS_ENTRY_PAYMENT'))

# Aggregate by current application ID
installments_agg = installments_df.groupBy(['SK_ID_CURR']).agg(
    F.count(col('SK_ID_PREV')).alias('NBR_SK_ID_PREV'),
    avg(col('AMT_LEFT')).alias('AVG_AMT_LEFT'),
    avg(col('DAYS_LATE')).alias('AVG_DAYS_LATE'),

(F.sum(col('AMT_LEFT'))/F.sum(col('AMT_INSTALLMENT'))).alias('PERC_AMT_
LEFT')
)

# Add prefix to column names for joining
installments_prefixed =
installments_agg.select([F.col(c).alias(f"IP_{c}") for c in
installments_agg.columns])

# Join with application_train
installments_merged = data['application_train'].join(
    installments_prefixed,
    data['application_train'].SK_ID_CURR ==
installments_prefixed.IP_SK_ID_CURR,
    how='inner'
)

# Analyze relationship with TARGET
installments_vars = [
    'IP_NBR_SK_ID_PREV',
    'IP_AVG_AMT_LEFT',
    'IP_AVG_DAYS_LATE',
    'IP_PERC_AMT_LEFT'
]
plot_dist_TARGET(installments_merged, installments_vars, num_rows=2)

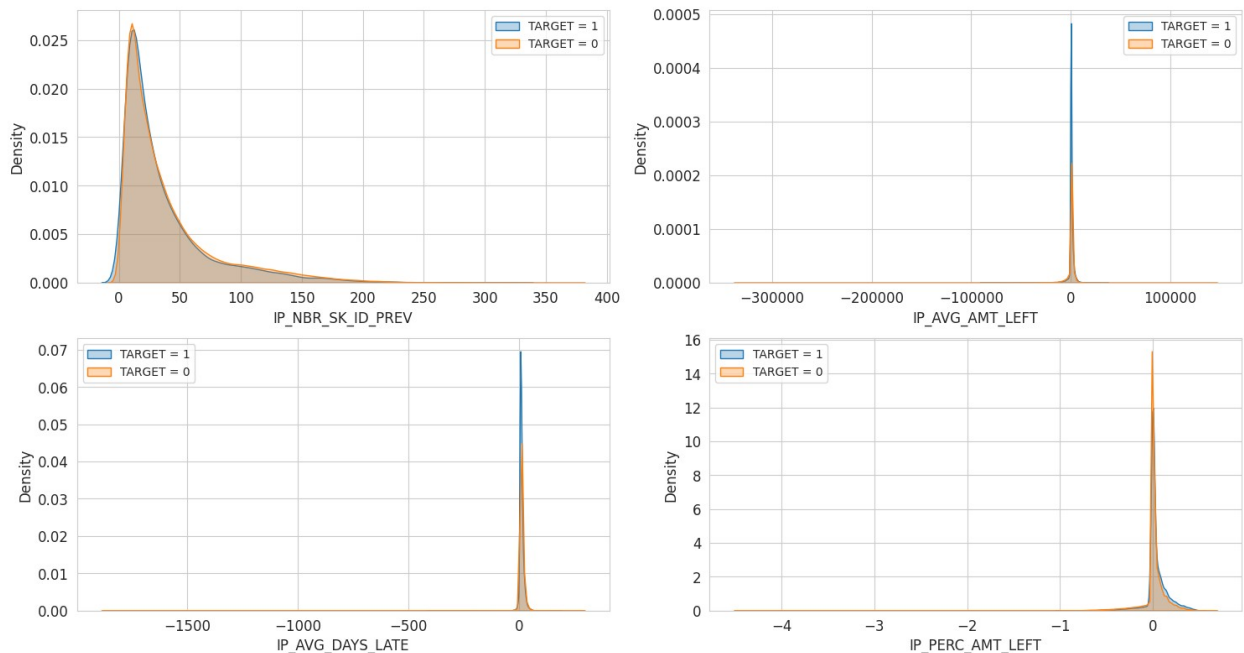
Correlation between IP_NBR_SK_ID_PREV and TARGET: -0.0211
Median value for default loans: 23.0000
Median value for non-default loans: 25.0000

```

Correlation between IP_AVG_AMT_LEFT and TARGET: 0.0293
Median value for default loans: 0.0000
Median value for non-default loans: 0.0000

Correlation between IP_AVG_DAYS_LATE and TARGET: -0.0209
Median value for default loans: 8.6316
Median value for non-default loans: 9.6250

Correlation between IP_PERC_AMT_LEFT and TARGET: 0.0527
Median value for default loans: 0.0000
Median value for non-default loans: 0.0000



2.2. Dataframe POS_CASH_balance

Monthly balance snapshots of previous POS (point of sales) and cash loans that the applicant had with Home Credit

```
key = 'POS_CASH_balance'  
print(f'Dataframe {key} includes {len(data[key].columns)} features and  
{data[key].count()} observations.')  
data[key].printSchema()
```

Dataframe POS_CASH_balance includes 8 features and 10001358 observations.

root

```
|-- SK_ID_PREV: integer (nullable = true)
|-- SK_ID_CURR: integer (nullable = true)
|-- MONTHS_BALANCE: integer (nullable = true)
|-- CNT_INSTALMENT: double (nullable = true)
|-- CNT_INSTALMENT_FUTURE: double (nullable = true)
|-- NAME_CONTRACT_STATUS: string (nullable = true)
|-- SK_DPD: integer (nullable = true)
|-- SK_DPD_DEF: integer (nullable = true)
```

Analyze contract status distribution

```
pos_df = data[key]
pos_df.groupBy("NAME_CONTRACT_STATUS").count()\
    .withColumn('percent', F.round(col('count')*100/pos_df.count(),
2))\
    .show()
```

NAME_CONTRACT_STATUS	count	percent
Demand	7065	0.07
Signed	87260	0.87
Active	9151119	91.5
Approved	4917	0.05
Completed	744883	7.45
Returned to the s...	5461	0.05
XNA	2	0.0
Canceled	15	0.0
Amortized debt	636	0.01

Aggregate by current application ID

```
pos_agg = pos_df.groupBy(['SK_ID_CURR']).agg(
    F.count(col('SK_ID_PREV')).alias('NBR_SK_ID_PREV'),
    avg(col('CNT_INSTALMENT_FUTURE')).alias('AVG_CNT_INSTALMENT_FUTURE'),
    avg(col('SK_DPD_DEF')).alias('AVG_SK_DPD_DEF'),
    (F.sum(col('CNT_INSTALMENT_FUTURE'))/F.sum(col('CNT_INSTALMENT'))).ali
as('PERC_CNT_LEFT')
)
```

Add prefix to column names for joining

```
pos_prefixed = pos_agg.select([F.col(c).alias(f"POS_{c}") for c in
pos_agg.columns])
```

Join with application_train

```

pos_merged = data['application_train'].join(
    pos_prefixed,
    data['application_train'].SK_ID_CURR ==
pos_prefixed.POS_SK_ID_CURR,
    how='inner'
)

# Analyze relationship with TARGET
pos_vars = [
    'POS_NBR_SK_ID_PREV',
    'POS_AVG_CNT_INSTALLMENT_FUTURE',
    'POS_AVG_SK_DPD_DEF',
    'POS_PERC_CNT_LEFT'
]
plot_dist_TARGET(pos_merged, pos_vars, num_rows=2)


Correlation between POS_NBR_SK_ID_PREV and TARGET: -0.0356
Median value for default loans: 19.0000
Median value for non-default loans: 23.0000

Correlation between POS_AVG_CNT_INSTALLMENT_FUTURE and TARGET: 0.0278
Median value for default loans: 7.2000
Median value for non-default loans: 6.9211

Correlation between POS_AVG_SK_DPD_DEF and TARGET: 0.0065
Median value for default loans: 0.0000
Median value for non-default loans: 0.0000

Correlation between POS_PERC_CNT_LEFT and TARGET: 0.0308
Median value for default loans: 0.5833
Median value for non-default loans: 0.5690

```



2.3. Dataframe credit_card_balance

```
key = 'credit_card_balance'
print(f'Dataframe {key} includes {len(data[key].columns)} features and {data[key].count()} observations.')
data[key].printSchema()
```

Dataframe credit_card_balance includes 23 features and 3840312 observations.

root

```
|-- SK_ID_PREV: integer (nullable = true)
|-- SK_ID_CURR: integer (nullable = true)
|-- MONTHS_BALANCE: integer (nullable = true)
|-- AMT_BALANCE: double (nullable = true)
|-- AMT_CREDIT_LIMIT_ACTUAL: integer (nullable = true)
|-- AMT_DRAWINGS_ATM_CURRENT: double (nullable = true)
|-- AMT_DRAWINGS_CURRENT: double (nullable = true)
|-- AMT_DRAWINGS_OTHER_CURRENT: double (nullable = true)
|-- AMT_DRAWINGS_POS_CURRENT: double (nullable = true)
|-- AMT_INST_MIN_REGULARITY: double (nullable = true)
|-- AMT_PAYMENT_CURRENT: double (nullable = true)
|-- AMT_PAYMENT_TOTAL_CURRENT: double (nullable = true)
|-- AMT_RECEIVABLE_PRINCIPAL: double (nullable = true)
|-- AMT_RECIVABLE: double (nullable = true)
|-- AMT_TOTAL_RECEIVABLE: double (nullable = true)
|-- CNT_DRAWINGS_ATM_CURRENT: double (nullable = true)
|-- CNT_DRAWINGS_CURRENT: integer (nullable = true)
|-- CNT_DRAWINGS_OTHER_CURRENT: double (nullable = true)
|-- CNT_DRAWINGS_POS_CURRENT: double (nullable = true)
|-- CNT_INSTALMENT_MATURE_CUM: double (nullable = true)
```

```
| -- NAME_CONTRACT_STATUS: string (nullable = true)
| -- SK_DPD: integer (nullable = true)
| -- SK_DPD_DEF: integer (nullable = true)
```

Analyze contract status distribution

```
cc_df = data[key]
cc_df.groupBy("NAME_CONTRACT_STATUS").count()\
    .withColumn('percent', F.round(col('count')*100/cc_df.count(),
2))\
    .show()
```

NAME_CONTRACT_STATUS	count	percent
Demand	1365	0.04
Sent proposal	513	0.01
Signed	11058	0.29
Refused	17	0.0
Active	3698436	96.31
Completed	128918	3.36
Approved	5	0.0

Aggregate by current application ID

```
cc_agg = cc_df.groupBy(['SK_ID_CURR']).agg(
    F.count(col('SK_ID_PREV')).alias('NBR_SK_ID_PREV'),

    avg(col('AMT_PAYMENT_TOTAL_CURRENT')).alias('AVG_AMT_PAYMENT_TOTAL'),
    avg(col('AMT_TOTAL_RECEIVABLE')).alias('AVG_SK_DPD_DEF'),

    F.sum(col('CNT_INSTALLMENT_MATURE_CUM')).alias('SUM_CNT_INSTALLMENT_MATURE')
)
```

Add prefix to column names for joining

```
cc_prefixed = cc_agg.select([F.col(c).alias(f"CR_{c}") for c in
cc_agg.columns])
```

Join with application_train

```
cc_merged = data['application_train'].join(
    cc_prefixed,
    data['application_train'].SK_ID_CURR == cc_prefixed.CR_SK_ID_CURR,
    how='inner'
)
```

Analyze relationship with TARGET

```
cc_vars = [
    'CR_NBR_SK_ID_PREV',
    'CR_AVG_AMT_PAYMENT_TOTAL',
```



```

'CR_AVG_SK_DPD_DEF',
'CR_SUM_CNT_INSTALMENT_MATURE',
]
plot_dist_TARGET(cc_merged, cc_vars, num_rows=2)

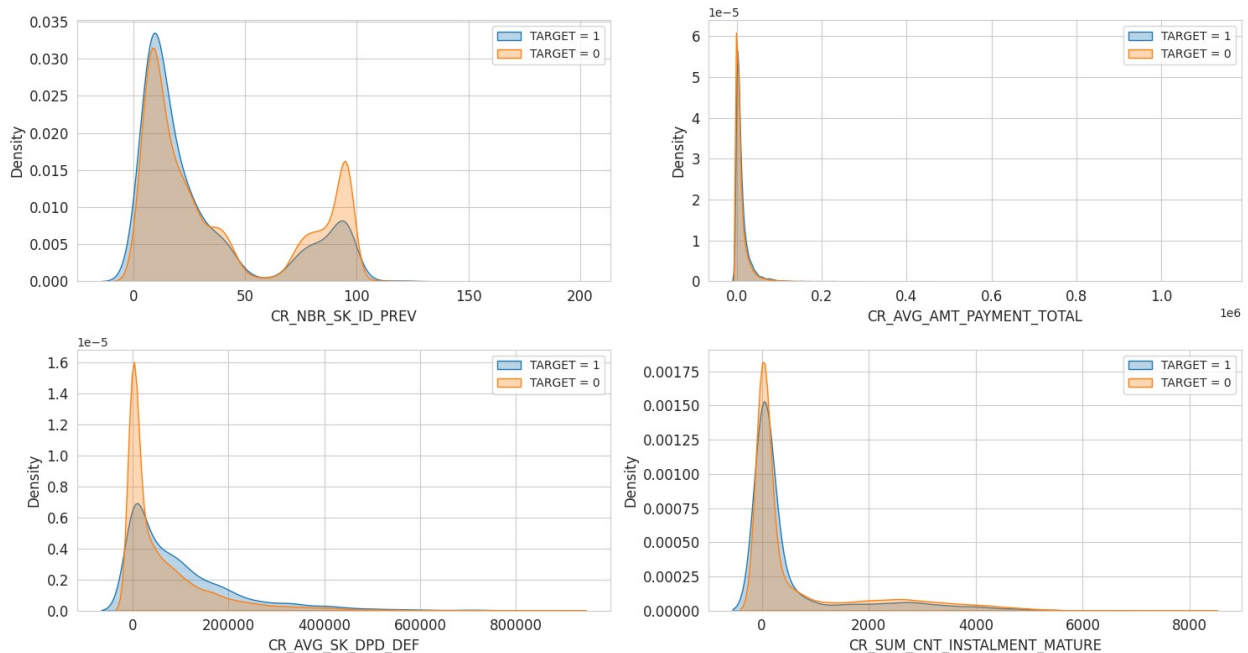
```

Correlation between CR_NBR_SK_ID_PREV and TARGET: -0.0605
Median value for default loans: 16.0000
Median value for non-default loans: 22.0000

Correlation between CR_AVG_AMT_PAYMENT_TOTAL and TARGET: 0.0227
Median value for default loans: 5985.5625
Median value for non-default loans: 3906.8182

Correlation between CR_AVG_SK_DPD_DEF and TARGET: 0.0865
Median value for default loans: 61815.1001
Median value for non-default loans: 23721.9764

Correlation between CR_SUM_CNT_INSTALMENT_MATURE and TARGET: -0.0424
Median value for default loans: 45.0000
Median value for non-default loans: 52.0000



2.4. Comment

Based on analysis of dataframes:

1. **bureau_balance:**
 - Contains monthly balances information from Credit Bureaus
 - Not strongly related to default rates in current applications
 - Will not be considered further
2. **installments_payments, POS_CASH_balance and credit_card_balance:**
 - Most applications have active status (>90%)
 - Refused/canceled applications are negligible (<1%)
 - Weak correlation with TARGET variable
 - No significant distributional differences between default/non-default groups

Therefore, we will focus mainly on **application_train/test, previous_application** and **bureau**.

3. EDA `application_train`

```
key = 'application_train'
print(f'Dataframe {key} includes {len(data[key].columns)} features and {data[key].count()} observations.')
data[key].printSchema()
```

Dataframe `application_train` includes 122 features and 307511 observations.

root

```
|-- SK_ID_CURR: integer (nullable = true)
|-- TARGET: integer (nullable = true)
|-- NAME_CONTRACT_TYPE: string (nullable = true)
|-- CODE_GENDER: string (nullable = true)
|-- FLAG_OWN_CAR: string (nullable = true)
|-- FLAG_OWN_REALTY: string (nullable = true)
|-- CNT_CHILDREN: integer (nullable = true)
|-- AMT_INCOME_TOTAL: double (nullable = true)
|-- AMT_CREDIT: double (nullable = true)
|-- AMT_ANNUITY: double (nullable = true)
|-- AMT_GOODS_PRICE: double (nullable = true)
|-- NAME_TYPE_SUITE: string (nullable = true)
|-- NAME_INCOME_TYPE: string (nullable = true)
|-- NAME_EDUCATION_TYPE: string (nullable = true)
|-- NAME_FAMILY_STATUS: string (nullable = true)
|-- NAME_HOUSING_TYPE: string (nullable = true)
|-- REGION_POPULATION_RELATIVE: double (nullable = true)
|-- DAYS_BIRTH: integer (nullable = true)
|-- DAYS_EMPLOYED: integer (nullable = true)
|-- DAYS_REGISTRATION: double (nullable = true)
|-- DAYS_ID_PUBLISH: integer (nullable = true)
|-- OWN_CAR_AGE: double (nullable = true)
|-- FLAG_MOBIL: integer (nullable = true)
|-- FLAG_EMP_PHONE: integer (nullable = true)
```

```
-- FLAG_WORK_PHONE: integer (nullable = true)
-- FLAG_CONT_MOBILE: integer (nullable = true)
-- FLAG_PHONE: integer (nullable = true)
-- FLAG_EMAIL: integer (nullable = true)
-- OCCUPATION_TYPE: string (nullable = true)
-- CNT_FAM_MEMBERS: double (nullable = true)
-- REGION_RATING_CLIENT: integer (nullable = true)
-- REGION_RATING_CLIENT_W_CITY: integer (nullable = true)
-- WEEKDAY_APPR_PROCESS_START: string (nullable = true)
-- HOUR_APPR_PROCESS_START: integer (nullable = true)
-- REG_REGION_NOT_LIVE_REGION: integer (nullable = true)
-- REG_REGION_NOT_WORK_REGION: integer (nullable = true)
-- LIVE_REGION_NOT_WORK_REGION: integer (nullable = true)
-- REG_CITY_NOT_LIVE_CITY: integer (nullable = true)
-- REG_CITY_NOT_WORK_CITY: integer (nullable = true)
-- LIVE_CITY_NOT_WORK_CITY: integer (nullable = true)
-- ORGANIZATION_TYPE: string (nullable = true)
-- EXT_SOURCE_1: double (nullable = true)
-- EXT_SOURCE_2: double (nullable = true)
-- EXT_SOURCE_3: double (nullable = true)
-- APARTMENTS_AVG: double (nullable = true)
-- BASEMENTAREA_AVG: double (nullable = true)
-- YEARS_BEGINEXPLUATATION_AVG: double (nullable = true)
-- YEARS_BUILD_AVG: double (nullable = true)
-- COMMONAREA_AVG: double (nullable = true)
-- ELEVATORS_AVG: double (nullable = true)
-- ENTRANCES_AVG: double (nullable = true)
-- FLOORSMAX_AVG: double (nullable = true)
-- FLOORSMIN_AVG: double (nullable = true)
-- LANDAREA_AVG: double (nullable = true)
-- LIVINGAPARTMENTS_AVG: double (nullable = true)
-- LIVINGAREA_AVG: double (nullable = true)
-- NONLIVINGAPARTMENTS_AVG: double (nullable = true)
-- NONLIVINGAREA_AVG: double (nullable = true)
-- APARTMENTS_MODE: double (nullable = true)
-- BASEMENTAREA_MODE: double (nullable = true)
-- YEARS_BEGINEXPLUATATION_MODE: double (nullable = true)
-- YEARS_BUILD_MODE: double (nullable = true)
-- COMMONAREA_MODE: double (nullable = true)
-- ELEVATORS_MODE: double (nullable = true)
-- ENTRANCES_MODE: double (nullable = true)
-- FLOORSMAX_MODE: double (nullable = true)
-- FLOORSMIN_MODE: double (nullable = true)
-- LANDAREA_MODE: double (nullable = true)
-- LIVINGAPARTMENTS_MODE: double (nullable = true)
-- LIVINGAREA_MODE: double (nullable = true)
-- NONLIVINGAPARTMENTS_MODE: double (nullable = true)
-- NONLIVINGAREA_MODE: double (nullable = true)
-- APARTMENTS_MEDI: double (nullable = true)
```

```
-- BASEMENTAREA_MEDI: double (nullable = true)
-- YEARS_BEGINEXPLUATATION_MEDI: double (nullable = true)
-- YEARS_BUILD_MEDI: double (nullable = true)
-- COMMONAREA_MEDI: double (nullable = true)
-- ELEVATORS_MEDI: double (nullable = true)
-- ENTRANCES_MEDI: double (nullable = true)
-- FLOORSMAX_MEDI: double (nullable = true)
-- FLOORSMIN_MEDI: double (nullable = true)
-- LANDAREA_MEDI: double (nullable = true)
-- LIVINGAPARTMENTS_MEDI: double (nullable = true)
-- LIVINGAREA_MEDI: double (nullable = true)
-- NONLIVINGAPARTMENTS_MEDI: double (nullable = true)
-- NONLIVINGAREA_MEDI: double (nullable = true)
-- FONDKAPREMONT_MODE: string (nullable = true)
-- HOUSETYPE_MODE: string (nullable = true)
-- TOTALAREA_MODE: double (nullable = true)
-- WALLSMATERIAL_MODE: string (nullable = true)
-- EMERGENCYSTATE_MODE: string (nullable = true)
-- OBS_30_CNT_SOCIAL_CIRCLE: double (nullable = true)
-- DEF_30_CNT_SOCIAL_CIRCLE: double (nullable = true)
-- OBS_60_CNT_SOCIAL_CIRCLE: double (nullable = true)
-- DEF_60_CNT_SOCIAL_CIRCLE: double (nullable = true)
-- DAYS_LAST_PHONE_CHANGE: double (nullable = true)
-- FLAG_DOCUMENT_2: integer (nullable = true)
-- FLAG_DOCUMENT_3: integer (nullable = true)
-- FLAG_DOCUMENT_4: integer (nullable = true)
-- FLAG_DOCUMENT_5: integer (nullable = true)
-- FLAG_DOCUMENT_6: integer (nullable = true)
-- FLAG_DOCUMENT_7: integer (nullable = true)
-- FLAG_DOCUMENT_8: integer (nullable = true)
-- FLAG_DOCUMENT_9: integer (nullable = true)
-- FLAG_DOCUMENT_10: integer (nullable = true)
-- FLAG_DOCUMENT_11: integer (nullable = true)
-- FLAG_DOCUMENT_12: integer (nullable = true)
-- FLAG_DOCUMENT_13: integer (nullable = true)
-- FLAG_DOCUMENT_14: integer (nullable = true)
-- FLAG_DOCUMENT_15: integer (nullable = true)
-- FLAG_DOCUMENT_16: integer (nullable = true)
-- FLAG_DOCUMENT_17: integer (nullable = true)
-- FLAG_DOCUMENT_18: integer (nullable = true)
-- FLAG_DOCUMENT_19: integer (nullable = true)
-- FLAG_DOCUMENT_20: integer (nullable = true)
-- FLAG_DOCUMENT_21: integer (nullable = true)
-- AMT_REQ_CREDIT_BUREAU_HOUR: double (nullable = true)
-- AMT_REQ_CREDIT_BUREAU_DAY: double (nullable = true)
-- AMT_REQ_CREDIT_BUREAU_WEEK: double (nullable = true)
-- AMT_REQ_CREDIT_BUREAU_MON: double (nullable = true)
-- AMT_REQ_CREDIT_BUREAU_QRT: double (nullable = true)
```

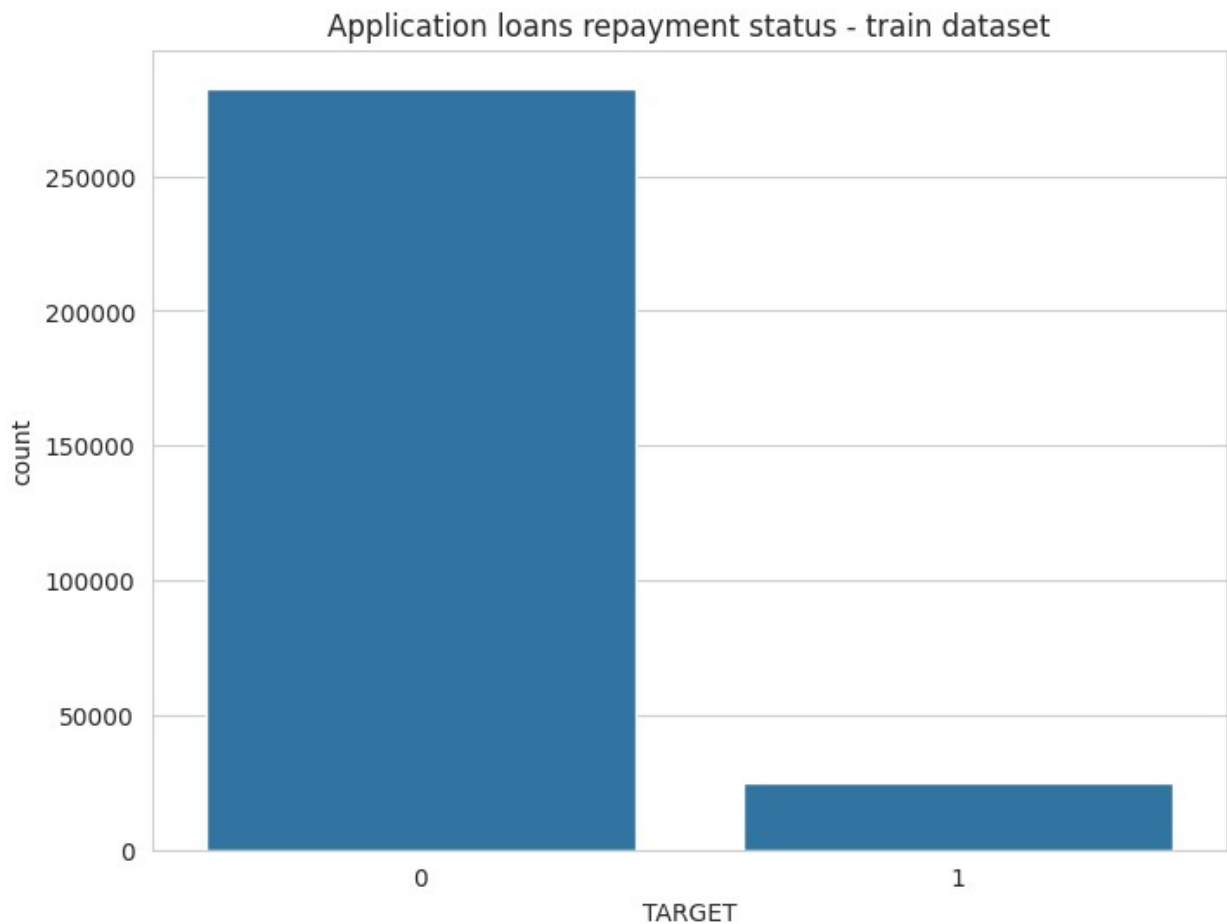
```
|-- AMT_REQ_CREDIT_BUREAU_YEAR: double (nullable = true)
```

3.1. Label TARGET

```
# Calculate age in years and employment duration
app_df = data[key]
app_df = app_df.withColumn('AGE_BIRTH', col('DAYS_BIRTH')/-365)
app_df = app_df.withColumn('YEARS_EMPLOYED', col('DAYS_EMPLOYED')/-365)

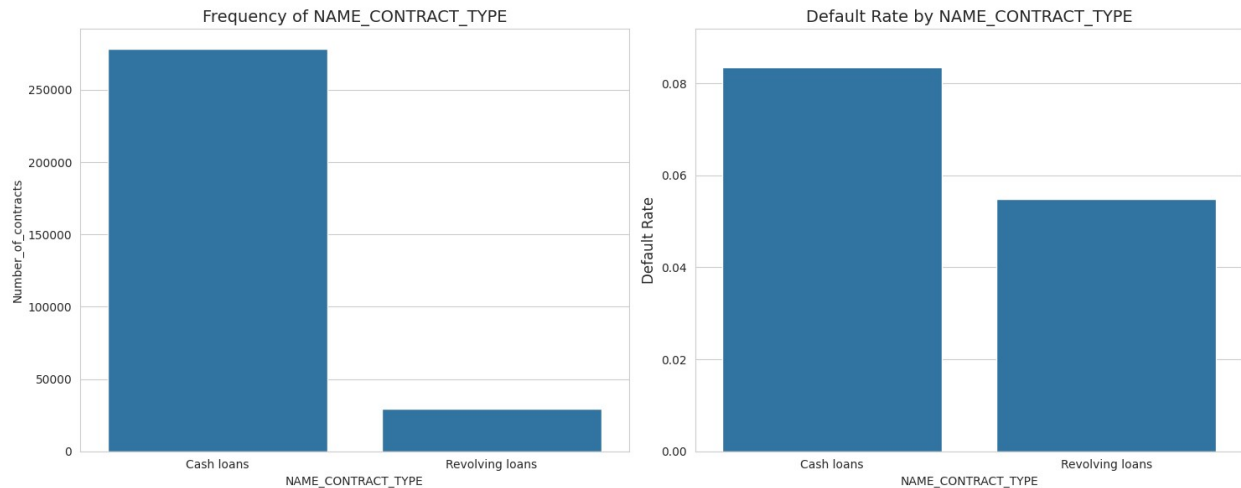
# Analyze target distribution
target_distribution = app_df.groupby('TARGET').count().toPandas()

# Plot target distribution
plt.figure(figsize=(8, 6))
plt.title('Application loans repayment status - train dataset')
sns.set_color_codes("pastel")
sns.barplot(x='TARGET', y="count", data=target_distribution)
locs, labels = plt.xticks()
plt.show()
```

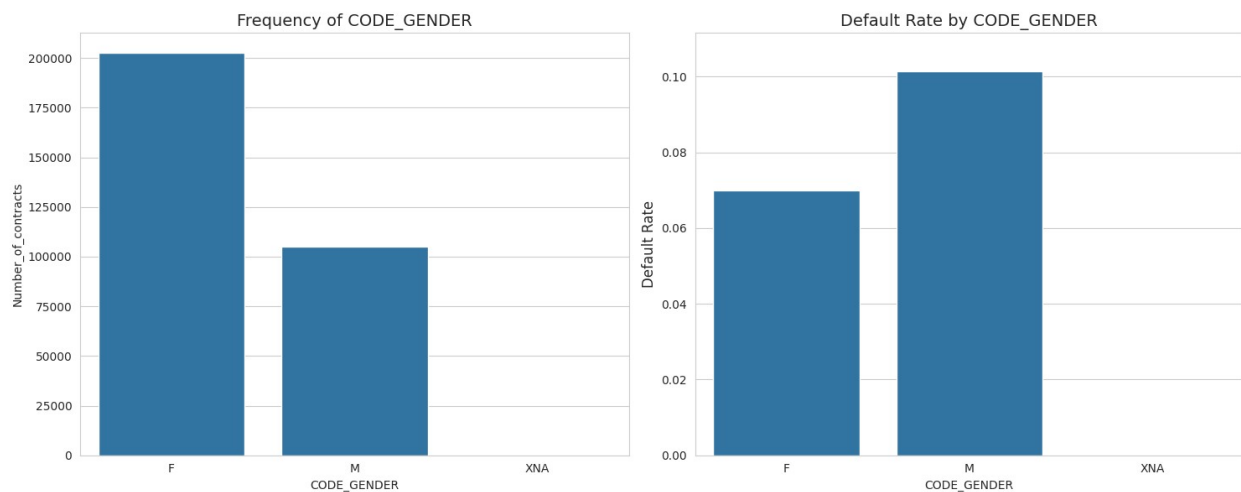


3.2. Categorical variables




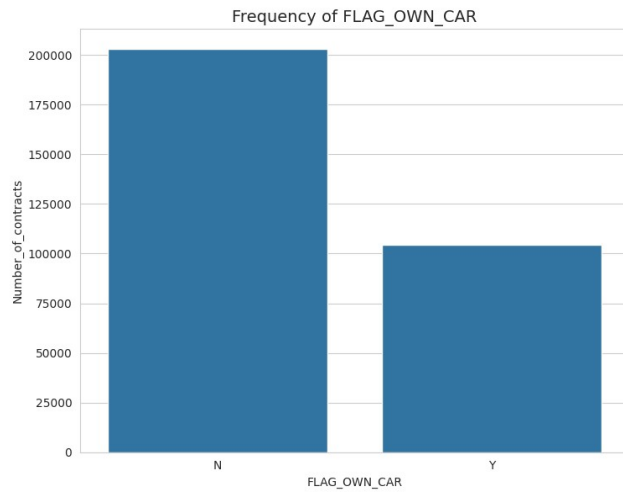
```
# Contract type  
plot_freq(app_df, 'NAME_CONTRACT_TYPE')
```



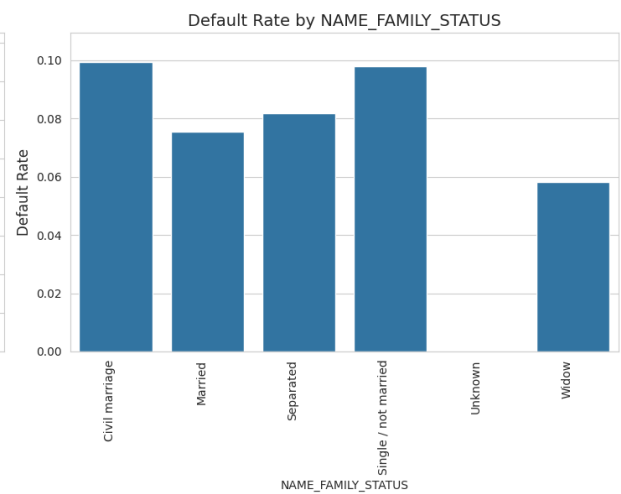
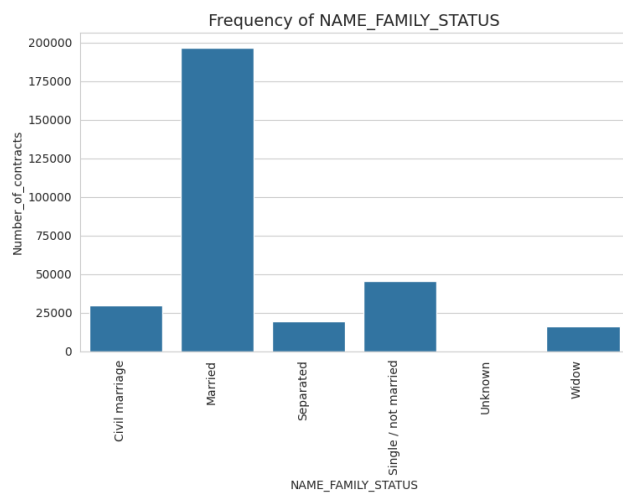
```
# Gender  
plot_freq(app_df, 'CODE_GENDER')
```



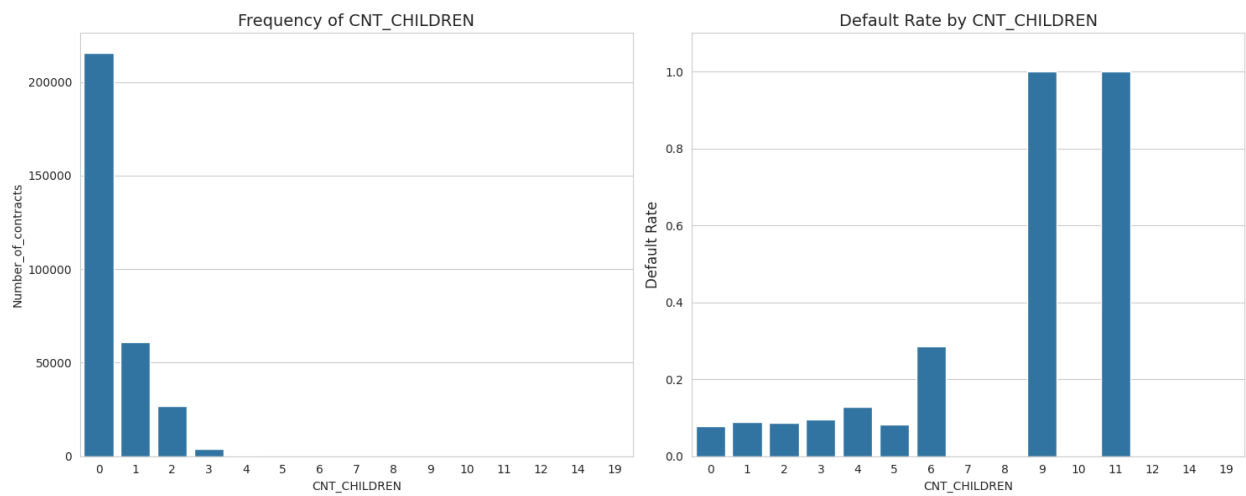
```
# Car and Property ownership  
plot_freq(app_df, 'FLAG_OWN_CAR')  
plot_freq(app_df, 'FLAG_OWN_REALTY')
```





```
# Family status
plot_freq(app_df, 'NAME_FAMILY_STATUS', True, True)
```



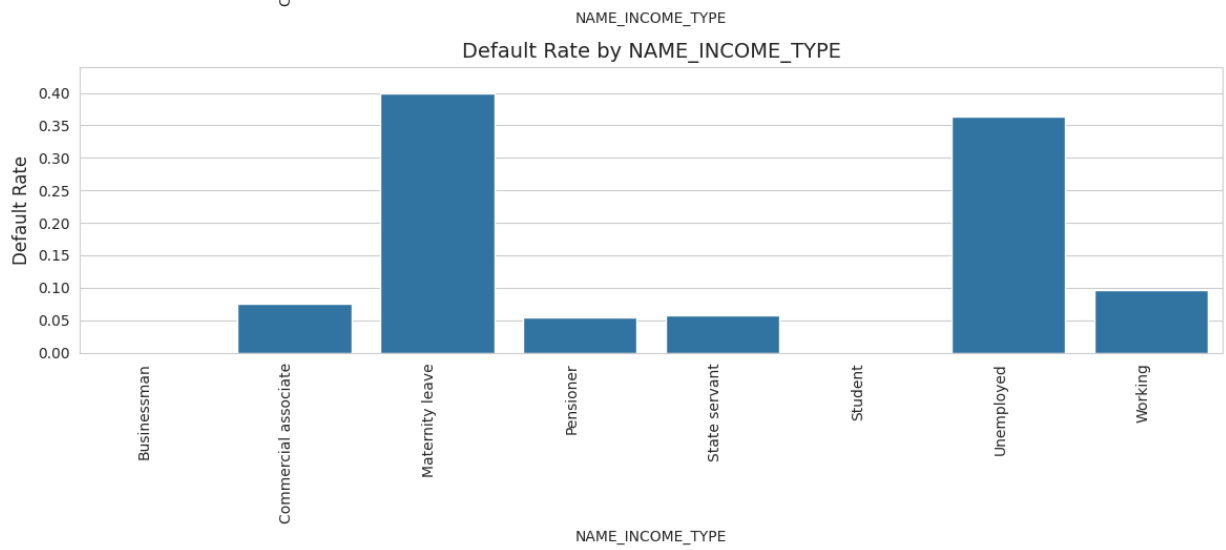
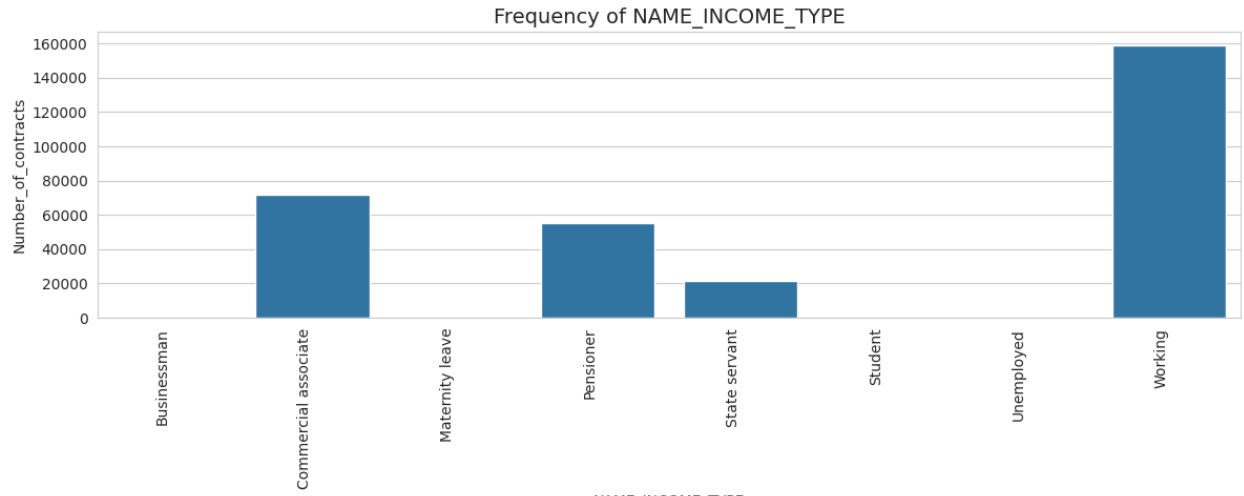
```
# Children count  
plot_freq(app_df, 'CNT_CHILDREN')
```





```
# Family members count  
plot_freq(app_df, 'CNT_FAM_MEMBERS', False, False)
```



```
# Income type  
plot_freq(app_df, 'NAME_INCOME_TYPE', False, True)
```



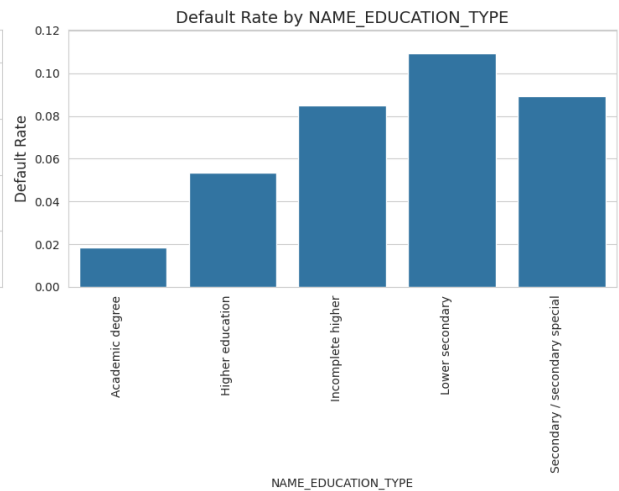
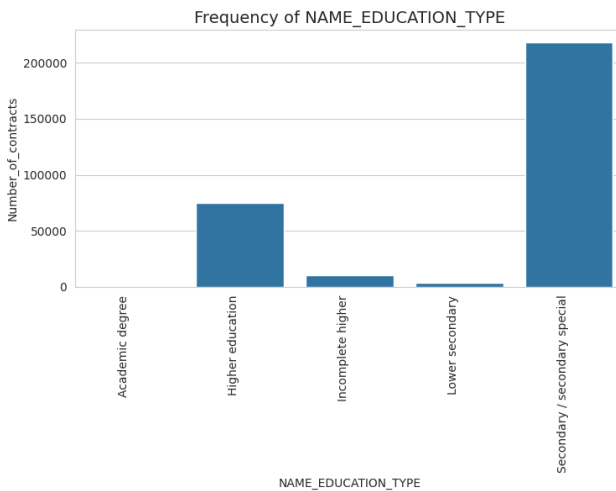
```
# Occupation type  
plot_freq(app_df, 'OCCUPATION_TYPE', False, True)
```



```
# Organization type
plot_freq(app_df, 'ORGANIZATION_TYPE', False, True)
```


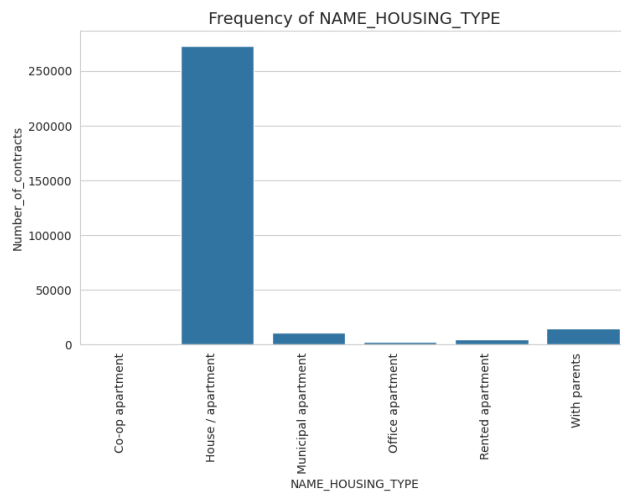


```
# Education type
plot_freq(app_df, 'NAME_EDUCATION_TYPE', True, True)
```



```
# Housing type
```

```
plot_freq(app_df, 'NAME_HOUSING_TYPE', True, True)
```




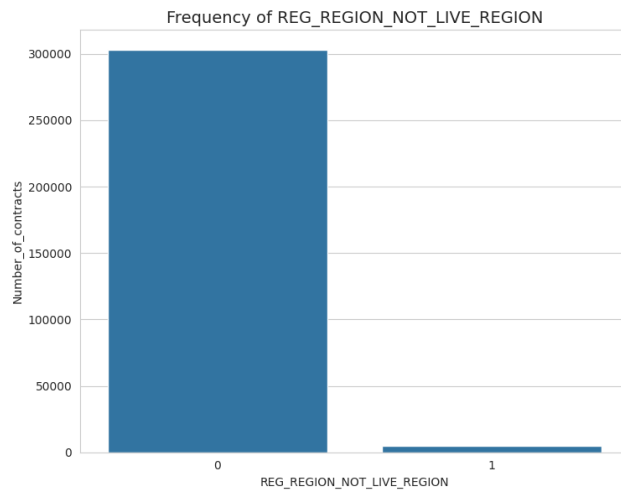
```
# Region flags
```


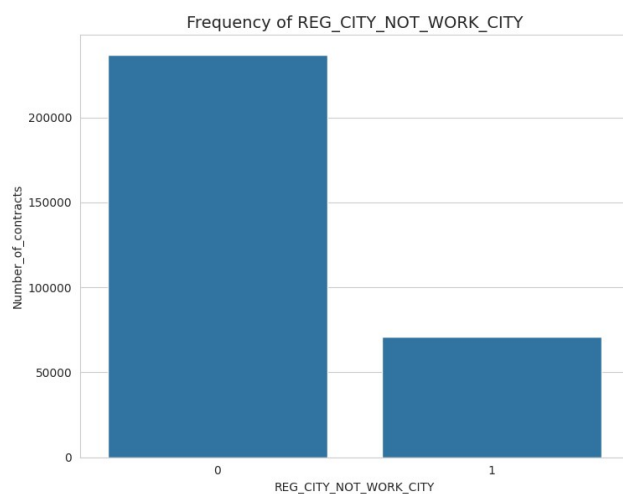

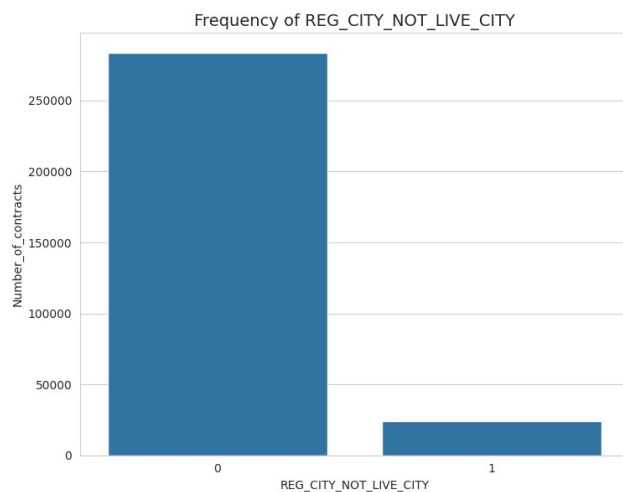
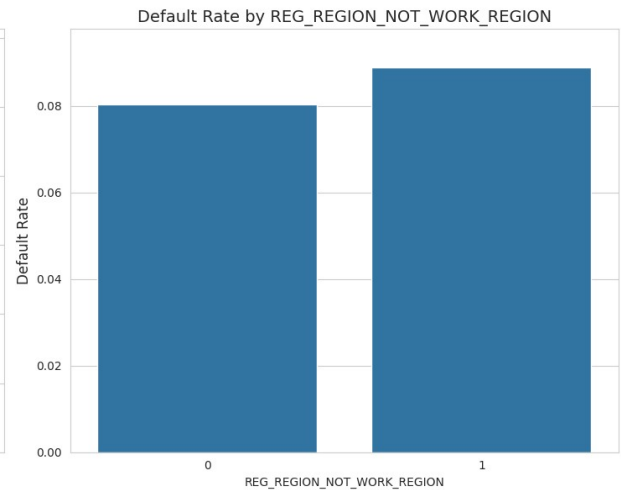

```
plot_freq(app_df, 'REG_REGION_NOT_LIVE_REGION')
```

```
plot_freq(app_df, 'REG_REGION_NOT_WORK_REGION')
```

```
plot_freq(app_df, 'REG_CITY_NOT_LIVE_CITY')
```


```
plot_freq(app_df, 'REG_CITY_NOT_WORK_CITY')
```






3.3. Numerical variables


```
# Income total
plot_dist(app_df, 'AMT_INCOME_TOTAL', 'red')
```




```
# Credit amount  
plot_dist(app_df, 'AMT_CREDIT', 'blue')
```




```
# Annuity amount  
plot_dist(app_df, 'AMT_ANNUITY', 'tomato')
```



```
# Goods price  
plot_dist(app_df, 'AMT_GOODS_PRICE', 'navy')
```




```
# Age  
plot_dist(app_df, 'AGE_BIRTH', 'green')
```


```
# Employment years  
plot_dist(app_df, 'YEARS_EMPLOYED', 'coral')
```



```
# Registration days  
plot_dist(app_df, 'DAYS_REGISTRATION', 'brown')
```



```
# Last phone change  
plot_dist(app_df, 'DAYS_LAST_PHONE_CHANGE', 'orange')
```



```
# Analyze numeric variables by TARGET  
numeric_vars = [  
    'AMT_ANNUITY',  
    'AMT_CREDIT',  
    'AMT_GOODS_PRICE',  
    'YEARS_EMPLOYED',  
    'DAYS_REGISTRATION',  
    'AGE_BIRTH',  
]
```

```
'EXT_SOURCE_1',  
'EXT_SOURCE_2',  
'EXT_SOURCE_3',  
'DAYS_LAST_PHONE_CHANGE'  
]  
plot_dist_TARGET(app_df, numeric_vars, num_rows=5)
```

Correlation between AMT_ANNUITY and TARGET: -0.0128
Median value for default loans: 25263.0000
Median value for non-default loans: 24876.0000

Correlation between AMT_CREDIT and TARGET: -0.0304
Median value for default loans: 497520.0000
Median value for non-default loans: 517788.0000

Correlation between AMT_GOODS_PRICE and TARGET: -0.0396
Median value for default loans: 450000.0000
Median value for non-default loans: 450000.0000

Correlation between YEARS_EMPLOYED and TARGET: 0.0449
Median value for default loans: 2.8329
Median value for non-default loans: 3.3836

Correlation between DAYS_REGISTRATION and TARGET: 0.0420
Median value for default loans: -4056.0000
Median value for non-default loans: -4544.0000

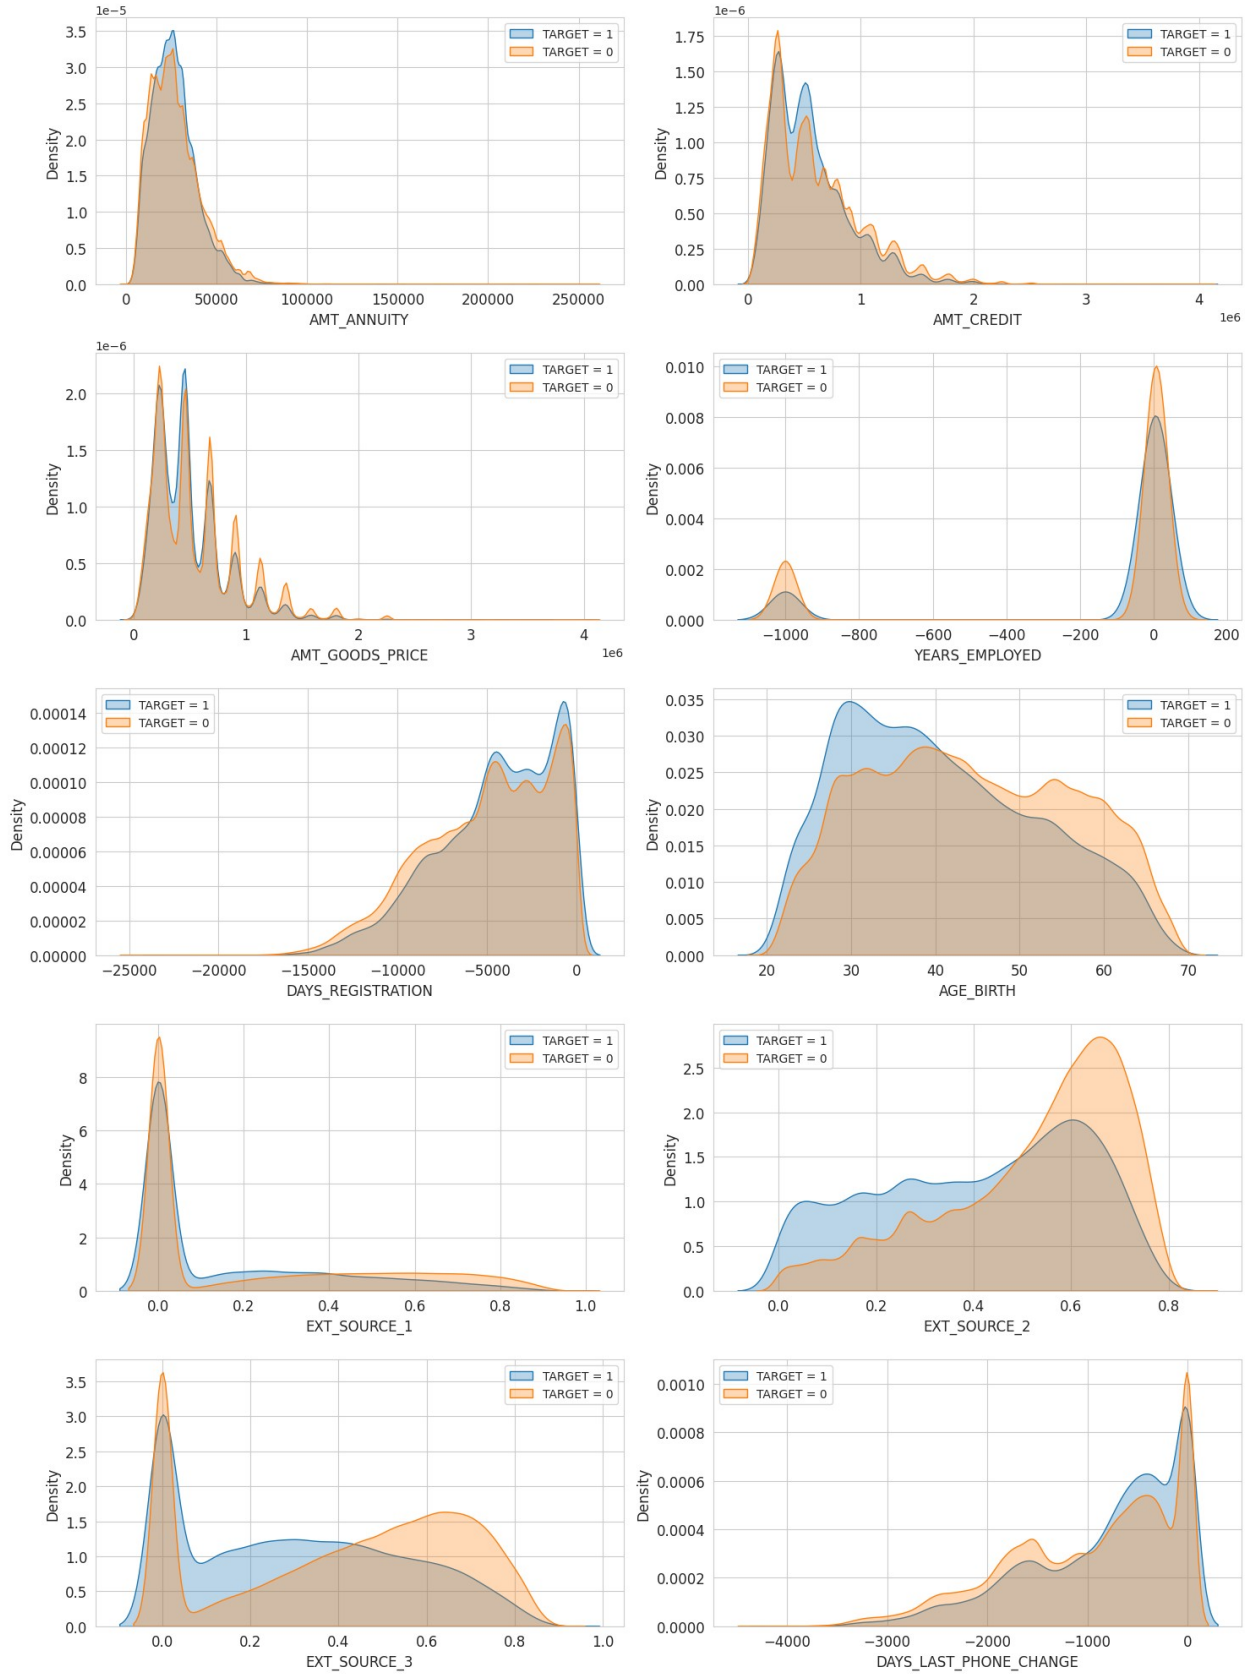
Correlation between AGE_BIRTH and TARGET: -0.0782
Median value for default loans: 39.1288
Median value for non-default loans: 43.4986

Correlation between EXT_SOURCE_1 and TARGET: -0.0647
Median value for default loans: 0.0000
Median value for non-default loans: 0.0000

Correlation between EXT_SOURCE_2 and TARGET: -0.1590
Median value for default loans: 0.4395
Median value for non-default loans: 0.5734

Correlation between EXT_SOURCE_3 and TARGET: -0.1196
Median value for default loans: 0.2881
Median value for non-default loans: 0.4741

Correlation between DAYS_LAST_PHONE_CHANGE and TARGET: 0.0552
Median value for default loans: -594.0000
Median value for non-default loans: -776.0000



4. EDA bureau

```
key = 'bureau'
print(f'Dataframe {key} includes {len(data[key].columns)} features and {data[key].count()} observations.')
data[key].printSchema()
```

Dataframe bureau includes 17 features and 1716428 observations.

```
root
|-- SK_ID_CURR: integer (nullable = true)
|-- SK_ID_BUREAU: integer (nullable = true)
|-- CREDIT_ACTIVE: string (nullable = true)
|-- CREDIT_CURRENCY: string (nullable = true)
|-- DAYS_CREDIT: integer (nullable = true)
|-- CREDIT_DAY_OVERDUE: integer (nullable = true)
|-- DAYS_CREDIT_ENDDATE: double (nullable = true)
|-- DAYS_ENDDATE_FACT: double (nullable = true)
|-- AMT_CREDIT_MAX_OVERDUE: double (nullable = true)
|-- CNT_CREDIT_PROLONG: integer (nullable = true)
|-- AMT_CREDIT_SUM: double (nullable = true)
|-- AMT_CREDIT_SUM_DEBT: double (nullable = true)
|-- AMT_CREDIT_SUM_LIMIT: double (nullable = true)
|-- AMT_CREDIT_SUM_OVERDUE: double (nullable = true)
|-- CREDIT_TYPE: string (nullable = true)
|-- DAYS_CREDIT_UPDATE: integer (nullable = true)
|-- AMT_ANNUITY: double (nullable = true)
```

Add prefix to bureau columns

```
bureau_prefixed = data['bureau'].select([F.col(c).alias(f"BUREAU_{c}")
for c in data['bureau'].columns])
```

Join with application_train

```
bureau_merged = data['application_train'].join(
    bureau_prefixed,
    data['application_train'].SK_ID_CURR ==
bureau_prefixed.BUREAU_SK_ID_CURR,
    how='left_outer'
)
```

Count total observations after join


```
print(f"Total observations after join: {bureau_merged.count()}")
```

Total observations after join: 1509345

4.1. Categorical features

Credit active status


```
plot_freq(bureau_merged, 'BUREAU_CREDIT_ACTIVE')
```



Observations:

- Most credits registered at Credit Bureau are Closed (~900K)
- Active credits account for ~600K
- Bad debt accounts for ~20% of defaults on current applications
- Clients with Closed credits have lowest default rate


```
# Credit currency
plot_freq(bureau_merged, 'BUREAU_CREDIT_CURRENCY')
```



Observations:

- Most credits are in currency_1
- Default rates vary by currency: currency_3 (11%), currency_1 (8%), currency_2 (5%)

```
# Credit type
plot_freq(bureau_merged, 'BUREAU_CREDIT_TYPE', True, True)
```




Observations:

- Majority are Consumer credit and Credit card
- Equipment purchase loans have >20% default rate
- Microloans have >20% default rate
- Working capital loans have >12% default rate

4.2. Numerical features

```
# Days credit
plot_dist(bureau_merged, 'BUREAU_DAYS_CREDIT', 'blue')
```




Observations:

- Credit duration ranges up to 3000 days
- Peak around 300 days (less than one year)


```
# Days overdue
bureau_merged.select('BUREAU_CREDIT_DAY_OVERDUE').toPandas().plot.hist
(bins=50)
```


```
<Axes: ylabel='Frequency'>
```



Observations:

- Most credits have 0 or very few days overdue
- Maximum overdue is ~3000 days

```
# Credit sum
plot_dist(bureau_merged, 'BUREAU_AMT_CREDIT_SUM', 'tomato')
```



Observations:

- Distribution shows several peaks
- Maximum concentration around 20,000

```
# Analyze bureau numeric variables by TARGET
```

```
bureau_numeric_vars = [
    'BUREAU_DAYS_CREDIT',
    'BUREAU_CREDIT_DAY_OVERDUE',
    'BUREAU_AMT_CREDIT_SUM',
    'BUREAU_AMT_CREDIT_SUM_LIMIT'
]
```

```
plot_dist_TARGET(bureau_merged, bureau_numeric_vars, num_rows=2)
```

Correlation between BUREAU_DAYS_CREDIT and TARGET: 0.0621

Median value for default loans: -726.0000

Median value for non-default loans: -970.0000

Correlation between BUREAU_CREDIT_DAY_OVERDUE and TARGET: 0.0025

Median value for default loans: 0.0000

Median value for non-default loans: 0.0000

Correlation between BUREAU_AMT_CREDIT_SUM and TARGET: -0.0112

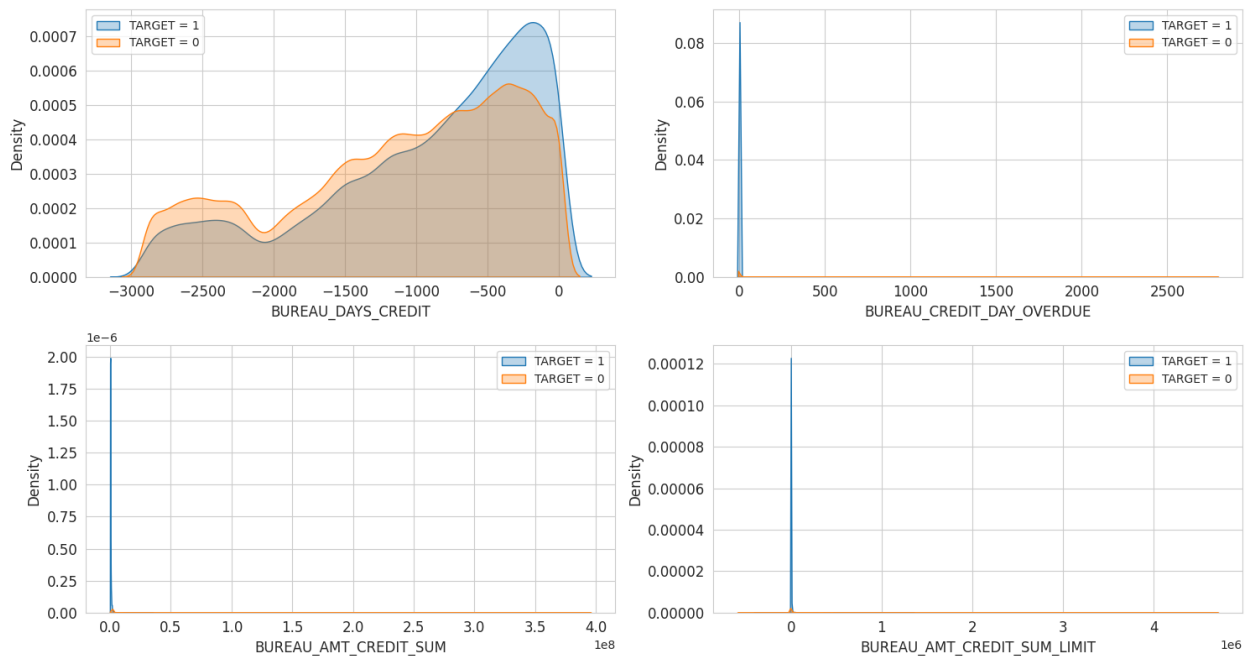
Median value for default loans: 112500.0000

Median value for non-default loans: 117000.0000

Correlation between BUREAU_AMT_CREDIT_SUM_LIMIT and TARGET: -0.0051

Median value for default loans: 0.0000

Median value for non-default loans: 0.0000



5. EDA `previous_application`

```
key = 'previous_application'
print(f'Dataframe {key} includes {len(data[key].columns)} features and {data[key].count()} observations.')
data[key].printSchema()
```

Dataframe `previous_application` includes 37 features and 1670214 observations.

root

```
-- SK_ID_PREV: integer (nullable = true)
-- SK_ID_CURR: integer (nullable = true)
-- NAME_CONTRACT_TYPE: string (nullable = true)
-- AMT_ANNUITY: double (nullable = true)
-- AMT_APPLICATION: double (nullable = true)
-- AMT_CREDIT: double (nullable = true)
-- AMT_DOWN_PAYMENT: double (nullable = true)
-- AMT_GOODS_PRICE: double (nullable = true)
-- WEEKDAY_APPR_PROCESS_START: string (nullable = true)
-- HOUR_APPR_PROCESS_START: integer (nullable = true)
-- FLAG_LAST_APPL_PER_CONTRACT: string (nullable = true)
-- NFLAG_LAST_APPL_IN_DAY: integer (nullable = true)
-- RATE_DOWN_PAYMENT: double (nullable = true)
-- RATE_INTEREST_PRIMARY: double (nullable = true)
-- RATE_INTEREST_PRIVILEGED: double (nullable = true)
```

```

|-- NAME_CASH_LOAN_PURPOSE: string (nullable = true)
|-- NAME_CONTRACT_STATUS: string (nullable = true)
|-- DAYS_DECISION: integer (nullable = true)
|-- NAME_PAYMENT_TYPE: string (nullable = true)
|-- CODE_REJECT_REASON: string (nullable = true)
|-- NAME_TYPE_SUITE: string (nullable = true)
|-- NAME_CLIENT_TYPE: string (nullable = true)
|-- NAME_GOODS_CATEGORY: string (nullable = true)
|-- NAME_PORTFOLIO: string (nullable = true)
|-- NAME_PRODUCT_TYPE: string (nullable = true)
|-- CHANNEL_TYPE: string (nullable = true)
|-- SELLERPLACE_AREA: integer (nullable = true)
|-- NAME_SELLER_INDUSTRY: string (nullable = true)
|-- CNT_PAYMENT: double (nullable = true)
|-- NAME_YIELD_GROUP: string (nullable = true)
|-- PRODUCT_COMBINATION: string (nullable = true)
|-- DAYS_FIRST_DRAWING: double (nullable = true)
|-- DAYS_FIRST_DUE: double (nullable = true)
|-- DAYS_LAST_DUE_1ST_VERSION: double (nullable = true)
|-- DAYS_LAST_DUE: double (nullable = true)
|-- DAYS_TERMINATION: double (nullable = true)
|-- NFLAG_INSURED_ON_APPROVAL: double (nullable = true)

# Add prefix to previous application columns
prev_prefixed = data[key].select([F.col(c).alias(f"PREV_{c}") for c in
data[key].columns])

# Join with application_train
prev_merged = data['application_train'].join(
    prev_prefixed,
    data['application_train'].SK_ID_CURR ==
prev_prefixed.PREV_SK_ID_CURR,
    how='left_outer'
)

# Count total observations after join
print(f"Total observations after join: {prev_merged.count()}")

Total observations after join: 1430155



```

5.1. Categorical variables

```

# Contract type
plot_freq(prev_merged, 'PREV_NAME_CONTRACT_TYPE', True, True)



```



Observations:

- Three types: Cash loans, Consumer loans, Revolving loans
- Cash and Consumer loans each ~600K, Revolving loans ~150K
- Default rates: Revolving loans (10%), Cash loans (9.5%), Consumer loans (8%)


```
# Cash loan purpose
plot_freq(prev_merged, 'PREV_NAME_CASH_LOAN_PURPOSE', True, True)
```



Observations:

- Main purposes: Repairs, Other, Urgent needs, Buying used car, Building a house
- Highest default rates: Refusal to name goal (23%), Hobby (20%), Car repairs (18%)


```
# Contract status
plot_freq(prev_merged, 'PREV_NAME_CONTRACT_STATUS', True, True)
```



Observations:

- Most are Approved (around 850K), followed by Canceled and Refused (around 240K)
- Highest default rates: Refused (12%), Canceled (9%), Unused offer (8%), Approved (<8%)



```
# Payment type
plot_freq(prev_merged, 'PREV_NAME_PAYMENT_TYPE', True, True)
```



Observations:

- Most paid via Cash through bank (~850K)
- All payment types show similar default rates (~8%)



```
# Client type
plot_freq(prev_merged, 'PREV_NAME_CLIENT_TYPE')
```



Observations:


- Most are Repeater (around 1M), New (around 200K), Refreshed (a round 100K)
- Default rates: New (8.5%), Repeater (8.25%), Refreshed (7%)

```
# Insurance flag
plot_freq(prev_merged, 'PREV_NFLAG_INSURED_ON_APPROVAL')
```




5.2. Numerical variables


```
# Annuity amount
plot_dist(prev_merged, 'PREV_AMT_ANNUIITY', 'green')
```



```
# Application amount  
plot_dist(prev_merged, 'PREV_AMT_APPLICATION', 'brown')
```



```
# Credit amount  
plot_dist(prev_merged, 'PREV_AMT_CREDIT', 'black')
```

```
# Analyze previous application numeric variables by TARGET
```

```
prev_numeric_vars = [  
    'PREV_AMT_ANNUITY',  
    'PREV_AMT_APPLICATION',  
    'PREV_AMT_CREDIT',  
    'PREV_DAYS_FIRST_DUE',  
    'PREV_DAYS_LAST_DUE',  
    'PREV_DAYS_TERMINATION'  
]  
plot_dist_TARGET(prev_merged, prev_numeric_vars, num_rows=3)
```

Correlation between PREV_AMT_ANNUITY and TARGET: -0.0197

Median value for default loans: 7307.8200

Median value for non-default loans: 8160.7950

Correlation between PREV_AMT_APPLICATION and TARGET: -0.0049

Median value for default loans: 65026.8000

Median value for non-default loans: 69525.0000

Correlation between PREV_AMT_CREDIT and TARGET: -0.0017

Median value for default loans: 73314.0000

Median value for non-default loans: 78826.5000

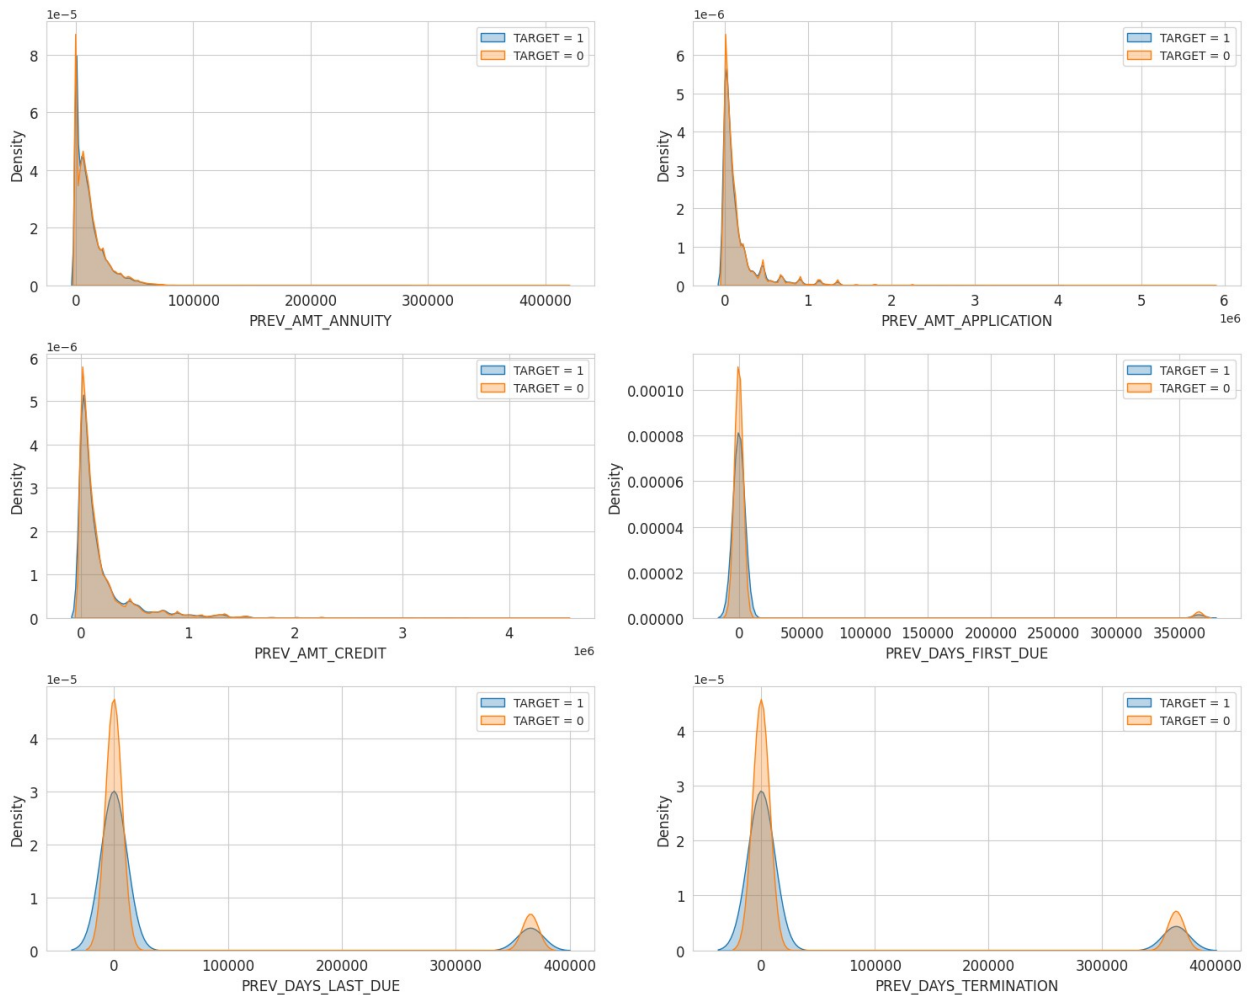
Correlation between PREV_DAYS_FIRST_DUE and TARGET: -0.0101

Median value for default loans: -56.0000

Median value for non-default loans: -288.0000

Correlation between PREV_DAYS_LAST_DUE and TARGET: -0.0017
Median value for default loans: 0.0000
Median value for non-default loans: 0.0000

Correlation between PREV_DAYS_TERMINATION and TARGET: -0.0028
Median value for default loans: 0.0000
Median value for non-default loans: 0.0000



6. Merging final dataset

```
# Select most important dataframes based on analysis
final_dataframes = [
    'application_train',
    'bureau',
    'previous_application'
```

```

]

# Select important columns from application_train
cols_application_train = [
    'SK_ID_CURR',
    'TARGET',
    'NAME_CONTRACT_TYPE',
    'CODE_GENDER',
    'FLAG_OWN_CAR',
    'FLAG_OWN_REALTY',
    'CNT_CHILDREN',
    'AMT_INCOME_TOTAL',
    'AMT_CREDIT',
    'AMT_ANNUITY',
    'AMT_GOODS_PRICE',
    'NAME_INCOME_TYPE',
    'NAME_EDUCATION_TYPE',
    'NAME_FAMILY_STATUS',
    'NAME_HOUSING_TYPE',
    'DAYS_BIRTH',
    'DAYS_EMPLOYED',
    'DAYS_REGISTRATION',
    'OCCUPATION_TYPE',
    'CNT_FAM_MEMBERS',
    'REG_REGION_NOT_LIVE_REGION',
    'REG_REGION_NOT_WORK_REGION',
    'REG_CITY_NOT_LIVE_CITY',
    'REG_CITY_NOT_WORK_CITY',
    'ORGANIZATION_TYPE',
    'EXT_SOURCE_1',
    'EXT_SOURCE_2',
    'EXT_SOURCE_3',
    'DAYS_LAST_PHONE_CHANGE'
]

# Select important columns from bureau
cols_bureau = [
    'SK_ID_CURR',
    'SK_ID_BUREAU',
    'CREDIT_ACTIVE',
    'DAYS_CREDIT',
    'DAYS_CREDIT_ENDDATE',
    'AMT_CREDIT_SUM',
    'AMT_CREDIT_SUM_OVERDUE',
    'CREDIT_TYPE'
]

# Select important columns from previous_application
cols_previous_application = [
    'SK_ID_PREV',

```

```

        'SK_ID_CURR',
        'NAME_CONTRACT_TYPE',
        'AMT_ANNUITY',
        'AMT_APPLICATION',
        'AMT_CREDIT',
        'AMT_GOODS_PRICE',
        'NAME_CONTRACT_STATUS',
        'NAME_CLIENT_TYPE',
        'DAYS_LAST_DUE',
        'DAYS_TERMINATION',
        'NFLAG_INSURED_ON_APPROVAL'
    ]

    # List of columns to keep for each dataframe
    columns_to_keep = [cols_application_train, cols_bureau,
                       cols_previous_application]

    # Dictionary to store filtered dataframes
    filtered_data = {}

    # Read and filter dataframes
    for idx, file in enumerate(final_dataframes):
        if file in filtered_data.keys():
            print(f'{file} already exists in dictionary')
        else:
            filepath = '/content/drive/MyDrive/Home_Credit/'+f'{file}.csv'
            filtered_df = spark.read.csv(
                filepath,
                header=True,
                sep=',',
                inferSchema=True
            ).select(columns_to_keep[idx])
            filtered_data[file] = filtered_df

    # Add prefix to bureau columns for joining
    bureau_prefixed =
    filtered_data['bureau'].select([F.col(c).alias(f"BUREAU_{c}") for c in
    filtered_data['bureau'].columns])

    # Add prefix to previous application columns for joining
    prev_app_prefixed =
    filtered_data['previous_application'].select([F.col(c).alias(f"PREV_{c}
    ") for c in filtered_data['previous_application'].columns])

    # Join application_train and bureau
    merged_df = filtered_data['application_train'].join(
        bureau_prefixed,
        filtered_data['application_train'].SK_ID_CURR ==
        bureau_prefixed.BUREAU_SK_ID_CURR,
        how='left_outer'
    )

```

```
)

# Join previous_application
final_df = merged_df.join(
    prev_app_prefixed,
    merged_df.SK_ID_CURR == prev_app_prefixed.PREV_SK_ID_CURR,
    how='left_outer'
)

# Count final observations
print(f"Final dataset has {final_df.count()} observations")

Final dataset has 8091522 observations

# Check number of distinct applicants
distinct_applicants =
final_df.select(F.countDistinct("SK_ID_CURR")).collect()[0][0]
print(f"Number of distinct applicants: {distinct_applicants}")

Number of distinct applicants: 307511

# Save final dataset to parquet file
final_df.write.mode("overwrite").parquet('/content/drive/MyDrive/Home_
Credit/train.parquet')

# Stop Spark session
spark.stop()
```